Yisheng Zhong

yzhong7@gmu.edu | easonzhong99.github.io | in/yisheng-zhong

Summary

Third-year Ph.D. student in Cybersecurity at **George Mason University** (advised by Dr. Zhuangdi Zhu), expected to graduate in 2028. My research focuses on **LLM post-training, safety alignment, unlearning, and agentic AI**.

Education

George Mason University, Ph.D. in Information Technology, GPA: 4.00 — 2024 – Present
University of Chinese Academy of Sciences, M.S. in Cyber Security, GPA: 3.56 — 2021 – 2024
Harbin University of Science and Technology, B.S. in Computer Science, GPA: 3.84 (Top 1%) — 2017 – 2021

Selected Publications

1. **DUET: Distilled LLM Unlearning from an Efficiently Contextualized Teacher** — *ICLR 2026*
   Yisheng Zhong, Zhengbang Yang, Zhuangdi Zhu
   - The first work to introduce **on-policy distillation** for LLM unlearning, where a prompt-conditioned teacher dynamically elicits desired forgetting behavior and distills it into the student model.
   - Introduced **Top-K logit alignment** for efficient and precise knowledge removal, achieving strong forget–retain trade-offs and improved robustness to reverse prompts and task-format shifts.
2. **Web Intellectual Property at Risk: Preventing Unauthorized Real-Time Retrieval by Large Language Models** — *EMNLP 2025, Main Conference*
   Yisheng Zhong, Yizhu Wen, Junfeng Guo, Mehran Kafai, Heng Huang, Hanqing Guo, Zhuangdi Zhu
   - Designed a **semantic defense framework** that embeds optimized HTML policy cues to prevent unauthorized real-time LLM retrieval, supporting refusal, partial masking, and source redirection.
   - Improved defense success rates from 2.5% to 88.6% across multiple proprietary LLMs, substantially outperforming configuration-based defenses such as robots.txt.
3. **CALIBURN: Self-Calibrated LLM Unlearning Alignment** — *EMNLP 2026, Main Conference*
   Zhengbang Yang, Yisheng Zhong, Junyuan Hong, Zhuangdi Zhu
   - Proposed **CALIBURN**, reformulating LLM unlearning as **policy-level preference optimization** with a self-calibrated margin derived from the model's own confidence, eliminating the need for a static reference model or retention data.
   - Developed **confidence-aware token-level gradient reweighting**, achieving state-of-the-art forget–retain trade-offs on WMDP and MUSE under scarce and heterogeneous unlearning data.

Industry Experience

The Washington Post | *Summer Intern, AI Team* — Summer 2026

- Developed an **LLM-assisted representation framework** for personalized news recommendation, introducing **StoryID** to capture the underlying story, entities, and article intent beyond article-level IDs or coarse topics.
- Integrated StoryID representations with existing recommendation models to improve personalization, with evaluation focused on **cold-start recommendation, content diversity, and recommendation interpretability**.

Honors & Awards

Meritorious Winner, Mathematical Contest in Modeling (**First Prize**), 2020